

COMP 3331/9331:
Computer Networks and
Applications

Week 6
Transport Layer (continued)

Reading Guide:
Chapter 3: Section 3.7

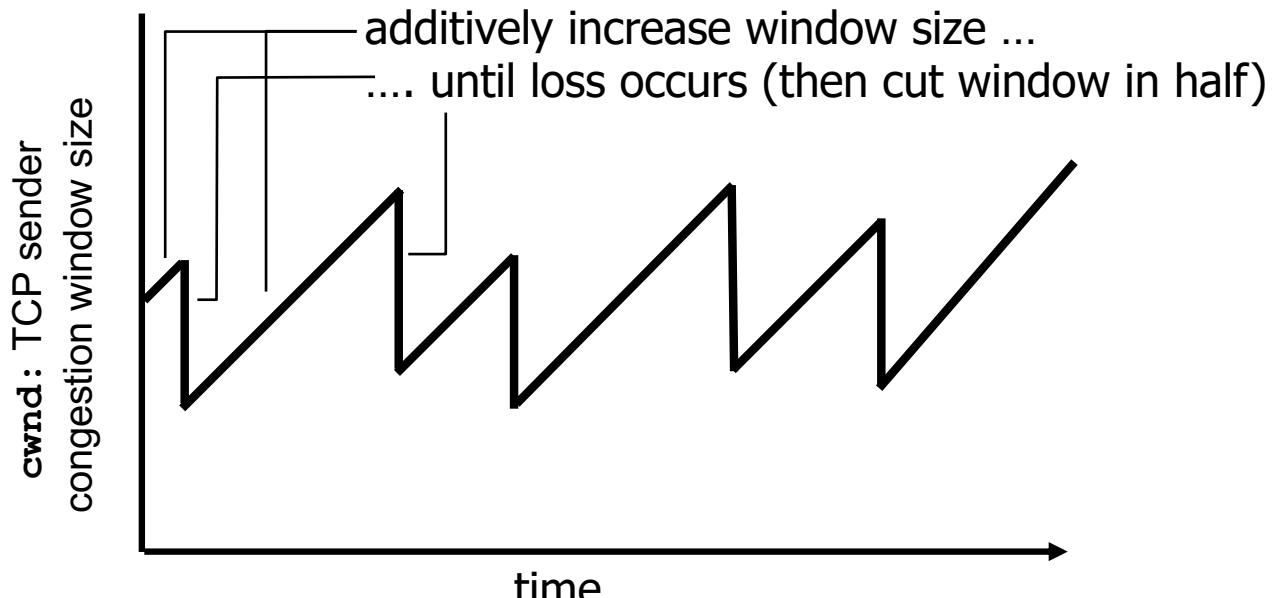
Adjusting to Varying Bandwidth

- Slow start gave an estimate of available bandwidth
- Now, want to track variations in this available bandwidth, oscillating around its current value
 - Repeated probing (rate increase) and backoff (rate decrease)
 - Known as Congestion Avoidance (CA)
- TCP uses: “Additive Increase Multiplicative Decrease” (AIMD)
 - We’ll see why shortly...

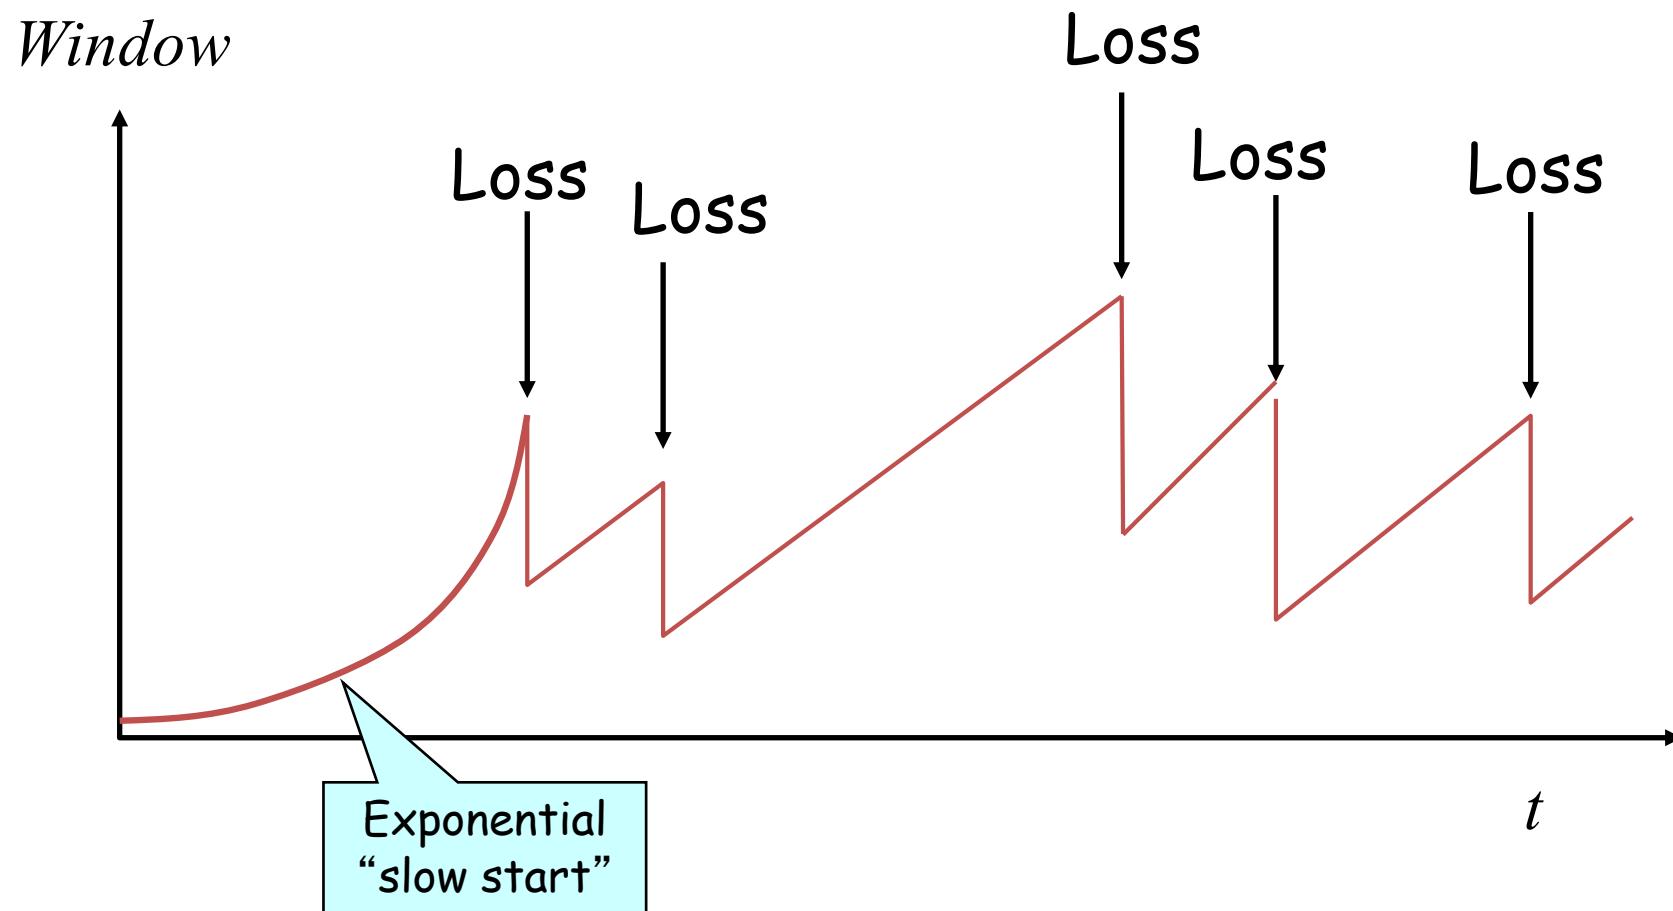
AIMD

- ❖ **approach:** sender increases transmission rate (window size), probing for usable bandwidth, until another congestion event occurs
 - **additive increase:** increase **cwnd** by 1 MSS every RTT until loss detected
 - For each successful RTT (all ACKS), $cwnd = cwnd + 1$
 - Simple implementation: for each ACK, $cwnd = cwnd + 1/cwnd$
 - **multiplicative decrease:** cut **cwnd** in half after loss

AIMD saw tooth behavior: probing for bandwidth



Leads to the TCP “Sawtooth”



Slow-Start vs. AIMD

- When does a sender stop Slow-Start and start Additive Increase?
- Introduce a “slow start threshold” (**ssthresh**)
 - Initialized to a large value
- Convert to AI when $cwnd = ssthresh$, sender switches from slow-start to AIMD-style increase
 - On timeout, $ssthresh = CWND/2$

Implementation

- State at sender
 - CWND (initialized to a small constant)
 - ssthresh (initialized to a large constant)
 - [Also dupACKcount and timer, as before]
- Events
 - ACK (new data)
 - dupACK (duplicate ACK for old data)
 - Timeout

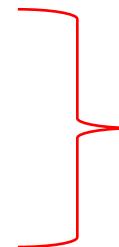
Event: ACK (new data)

- If $CWND < ssthresh$
 - $CWND += 1$

- $2 * MSS$ packets per ACK
- Hence after one RTT (All ACKs with no drops):
 $CWND = 2 \times CWND$

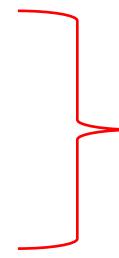
Event: ACK (new data)

- If $CWND < ssthresh$
 - $CWND += 1$



Slow start phase

- Else
 - $CWND = \min(CWND + 1, CWND / CWND)$



*“Congestion
Avoidance” phase
(additive increase)*

- *Hence after one RTT (All ACKs with no drops):*
 $CWND = CWND + 1$

Event: dupACK

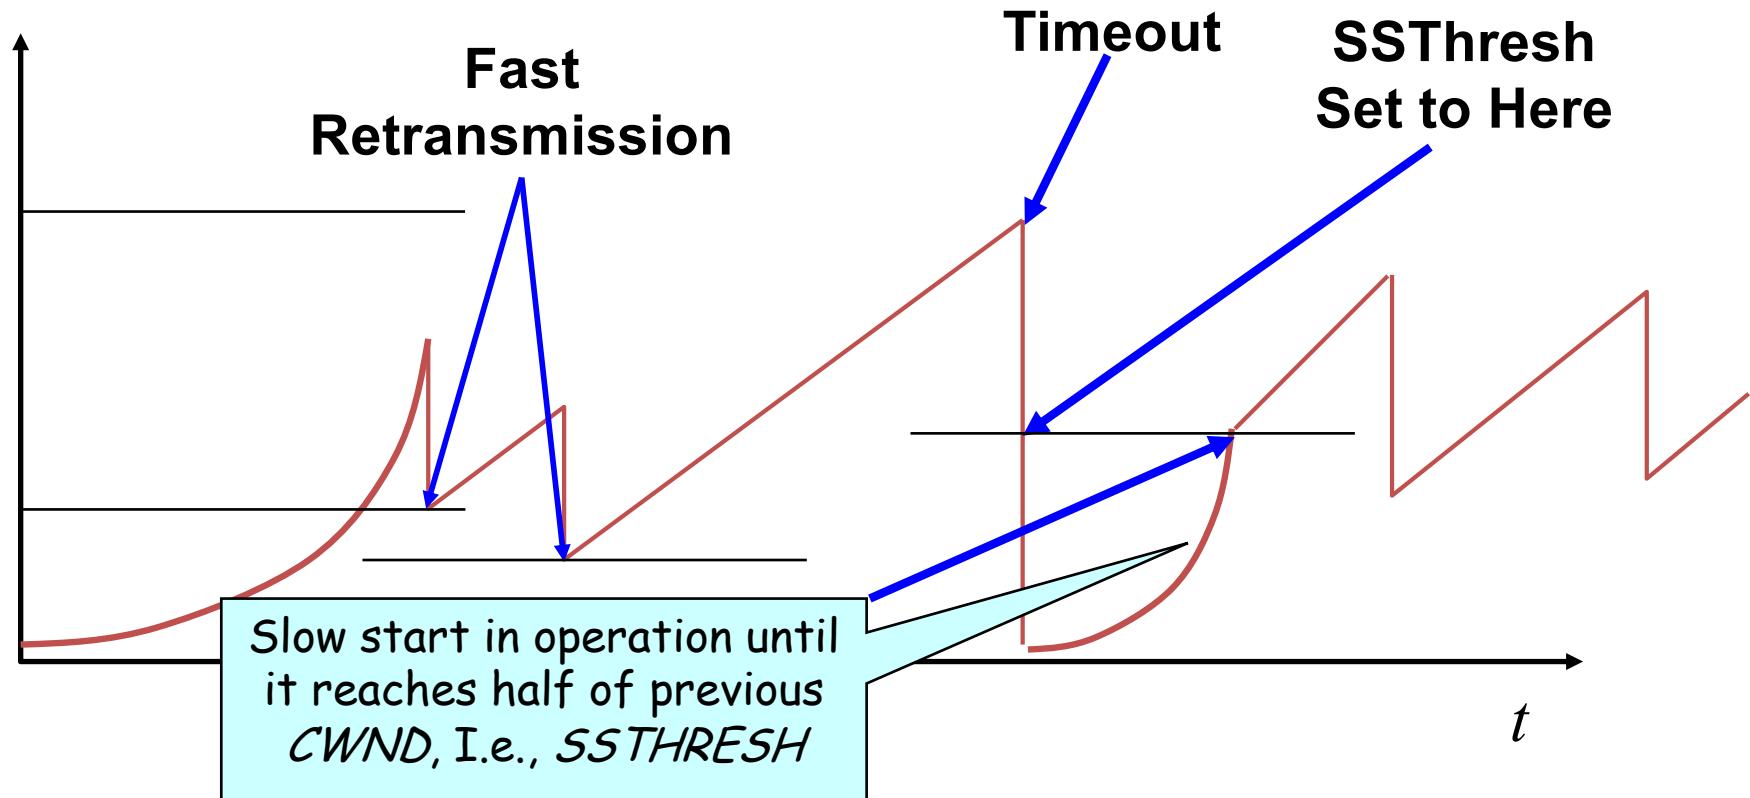
- dupACKcount ++
- If dupACKcount = 3 /* fast retransmit */
 - ssthresh = CWND/2
 - CWND = CWND/2

Event: TimeOut

- On Timeout
 - $ssthresh = CWND/2$
 - $CWND = 1$

Example

Window



Slow-start restart: Go back to $CWND = 1$ MSS, but take advantage of knowing the previous value of $CWND$

One Final Phase: Fast Recovery (improved)

- The problem: Fast retransmit slow in recovering from an isolated loss

Example (window in units of MSS, not bytes)

- Consider a TCP connection with:
 - CWND=10 packets (of size MSS, which is 100 bytes)
 - Last ACK was for byte # 101
 - i.e., receiver expecting next packet to have seq. no. 101
- 10 packets [101, 201, 301,..., 1001] are in flight
 - Packet 101 is dropped
 - What ACKs do they generate?
 - And how does the sender respond?

Timeline

- ACK 101 (due to 201) cwnd=10 dupACK#1 (no xmit)
- ACK 101 (due to 301) cwnd=10 dupACK#2 (no xmit)
- ACK 101 (due to 401) cwnd=10 dupACK#3 (no xmit)
- RETRANSMIT 101 ssthresh=5 cwnd= 5
- ACK 101 (due to 501) cwnd=5 + 1/5 (no xmit)
- ACK 101 (due to 601) cwnd=5 + 2/5 (no xmit)
- ACK 101 (due to 701) cwnd=5 + 3/5 (no xmit)
- ACK 101 (due to 801) cwnd=5 + 4/5 (no xmit)
- ACK 101 (due to 901) cwnd=5 + 5/5 (no xmit)
- ACK 101 (due to 1001) cwnd=6 + 1/6 (no xmit)
- ACK 1101 (due to 101) ← only now can we transmit new packets
- Plus no packets in flight so ACK “clocking” (to increase CWND) stalls for another RTT

Solution: Fast Recovery (improved)

Idea: Grant the sender temporary “credit” for each dupACK so as to keep packets in flight

- If $\text{dupACKcount} = 3$
 - $\text{ssthresh} = \text{cwnd}/2$
 - $\text{cwnd} = \text{ssthresh} + 3$
- While in fast recovery
 - $\text{cwnd} = \text{cwnd} + 1$ for each additional duplicate ACK
- Exit fast recovery after receiving new ACK
 - set $\text{cwnd} = \text{ssthresh}$

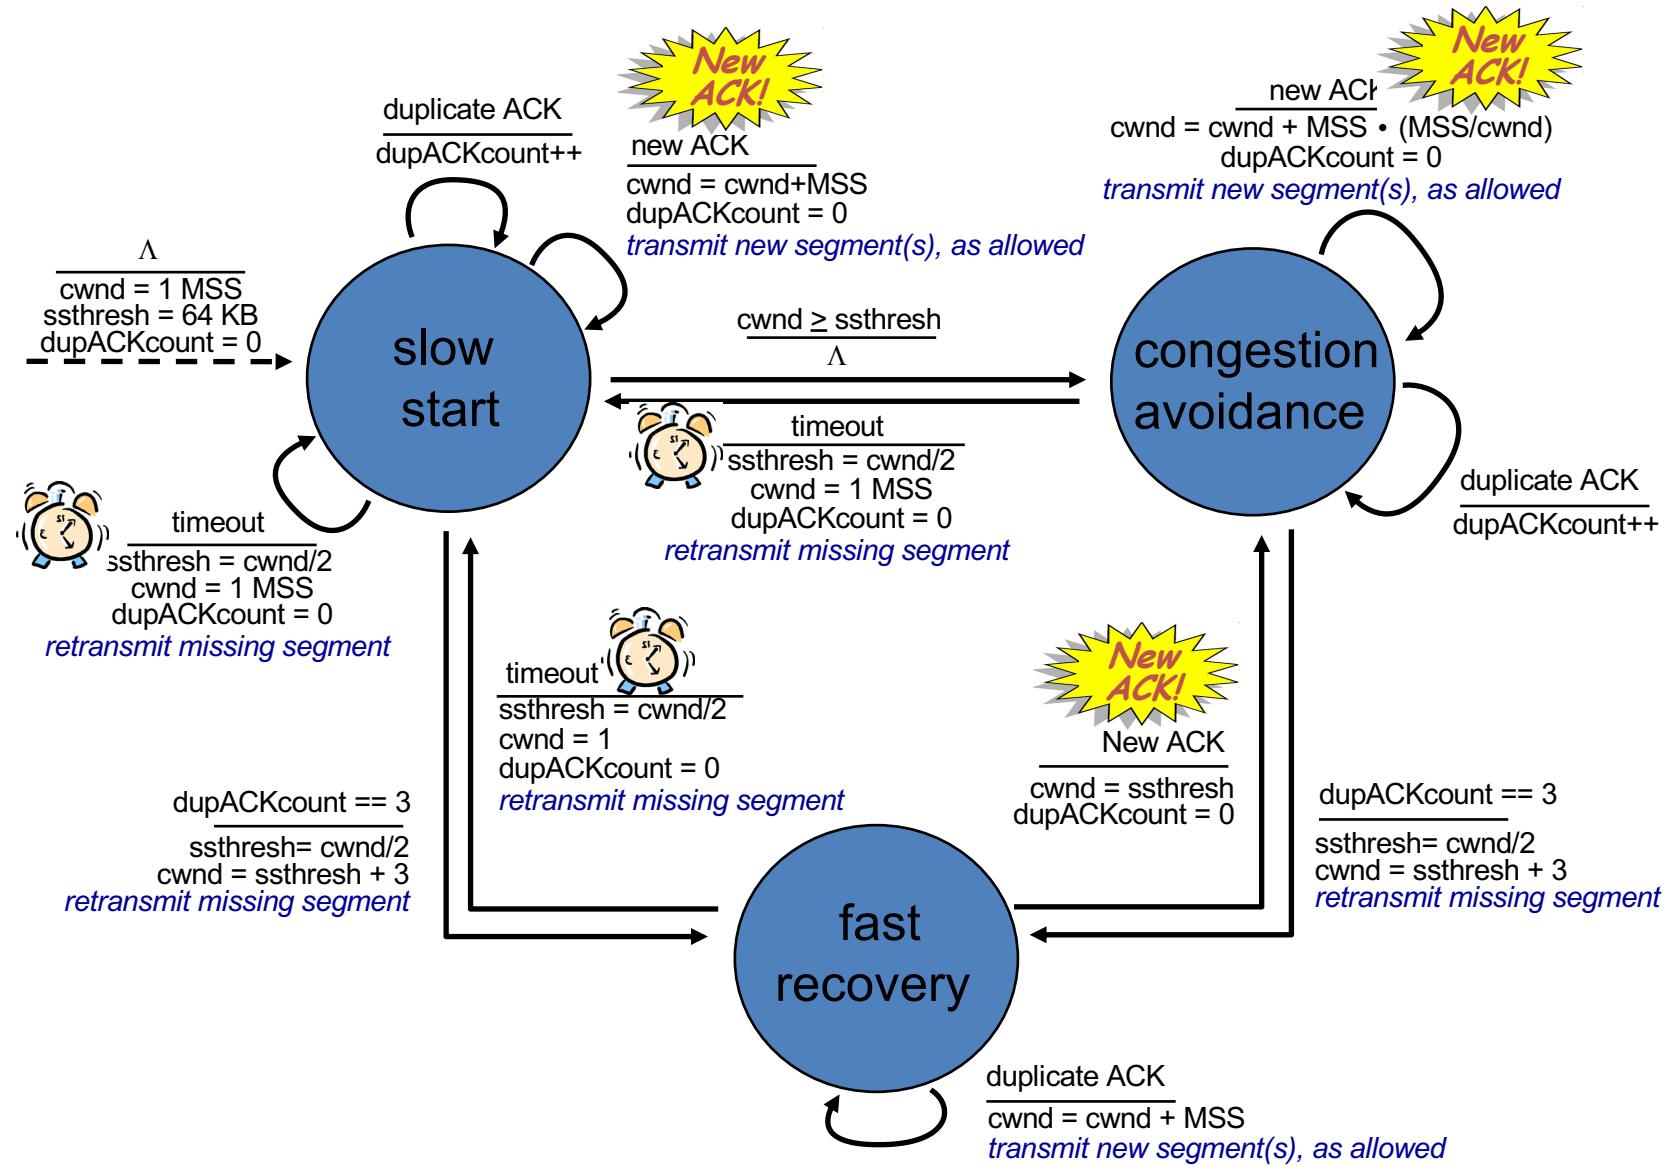
Example

- Consider a TCP connection with:
 - CWND=10 packets (of size MSS = 100 bytes)
 - Last ACK was for byte # 101
 - i.e., receiver expecting next packet to have seq. no. 101
- 10 packets [101, 201, 301,..., 1001] are in flight
 - Packet 101 is dropped

Timeline

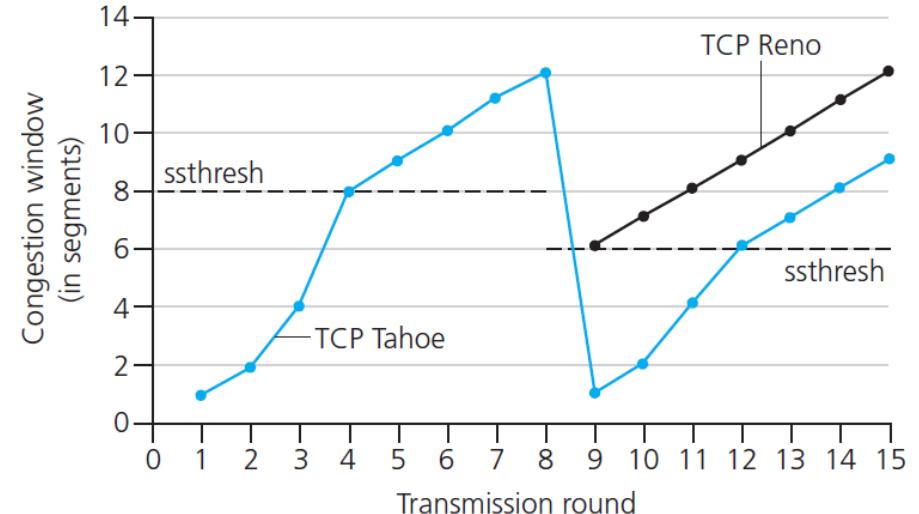
- ACK 101 (due to 201) cwnd=10 dup#1
- ACK 101 (due to 301) cwnd=10 dup#2
- ACK 101 (due to 401) cwnd=10 dup#3
- REXMIT 101 ssthresh=5 cwnd= 8 (5+3)
- ACK 101 (due to 501) cwnd= 9 (no xmit)
- ACK 101 (due to 601) cwnd=10 (no xmit)
- ACK 101 (due to 701) cwnd=11 (xmit 1101)
- ACK 101 (due to 801) cwnd=12 (xmit 1201)
- ACK 101 (due to 901) cwnd=13 (xmit 1301)
- ACK 101 (due to 1001) cwnd=14 (xmit 1401)
- ACK 1101 (due to 101) cwnd = 5 (xmit 1501) ← exiting fast recovery
- Packets 1101-1401 already in flight
- ACK 1201 (due to 1101) cwnd = 5 + 1/5 ← back in congestion avoidance

Summary: TCP Congestion Control



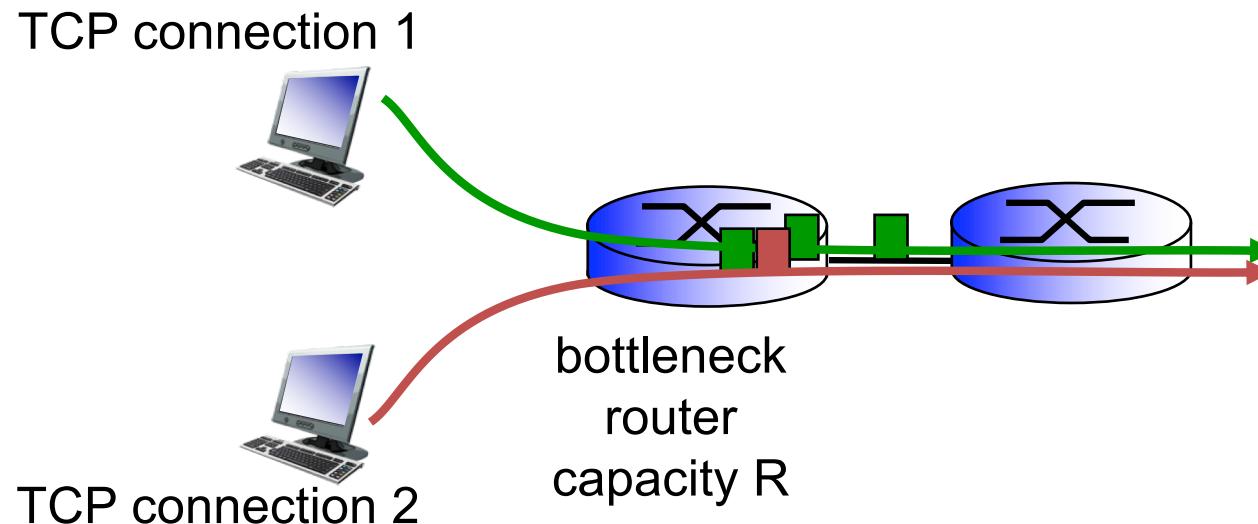
TCP Flavours

- TCP-Tahoe
 - cwnd = 1 on triple dup ACK & timeout
- TCP-Reno
 - cwnd = 1 on timeout
 - cwnd = cwnd/2 on triple dup ACK
- TCP-newReno
 - TCP-Reno + improved fast recovery
- TCP-SACK (NOT COVERED IN THE COURSE)
 - incorporates selective acknowledgements



TCP Fairness

fairness goal: if K TCP sessions share same bottleneck link of bandwidth R , each should have average rate of R/K

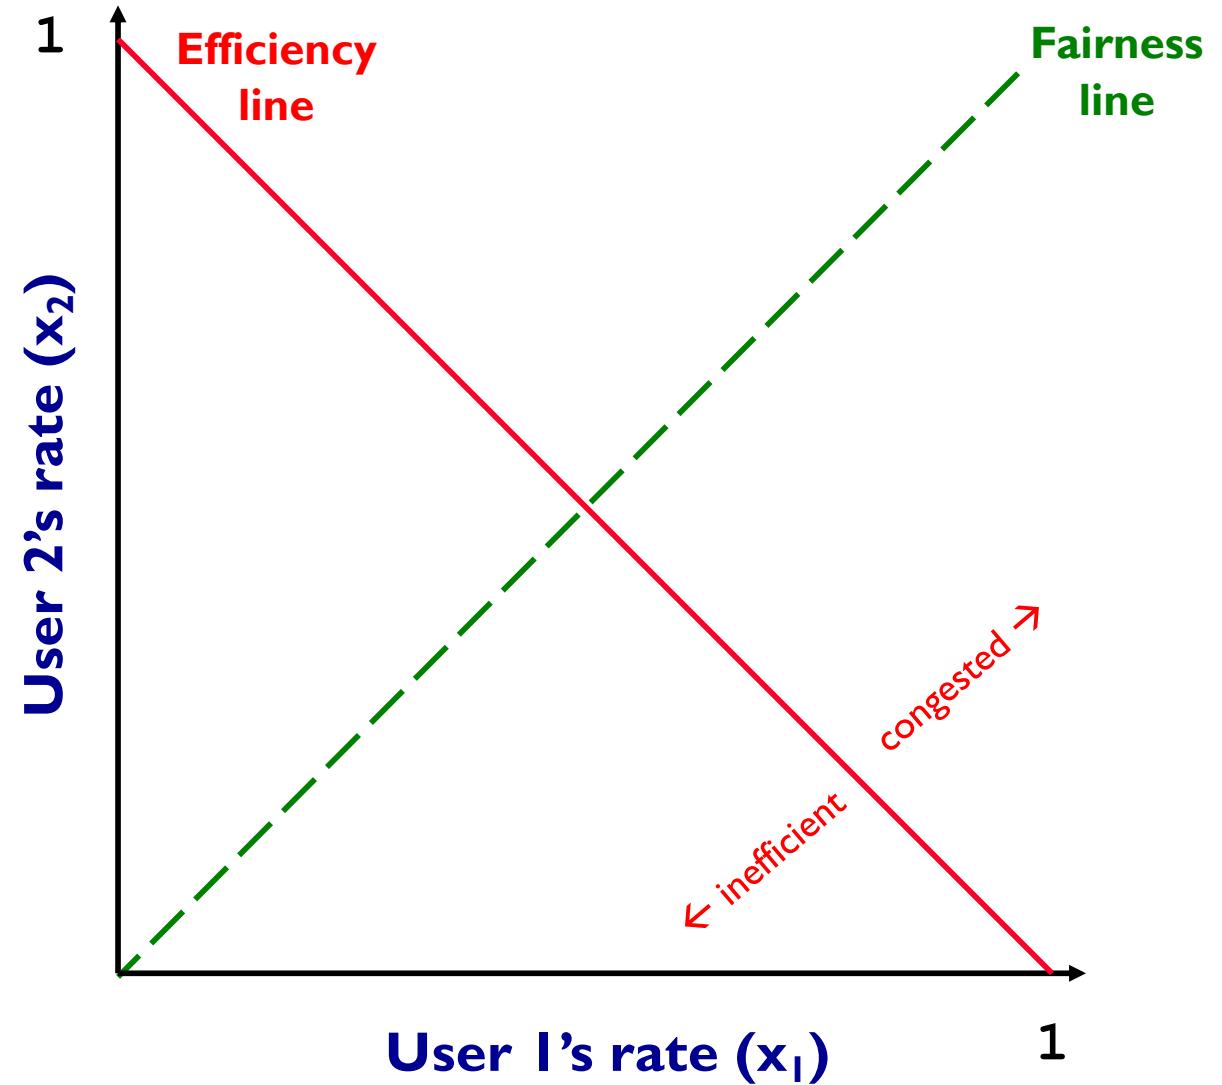


Why AIMD?

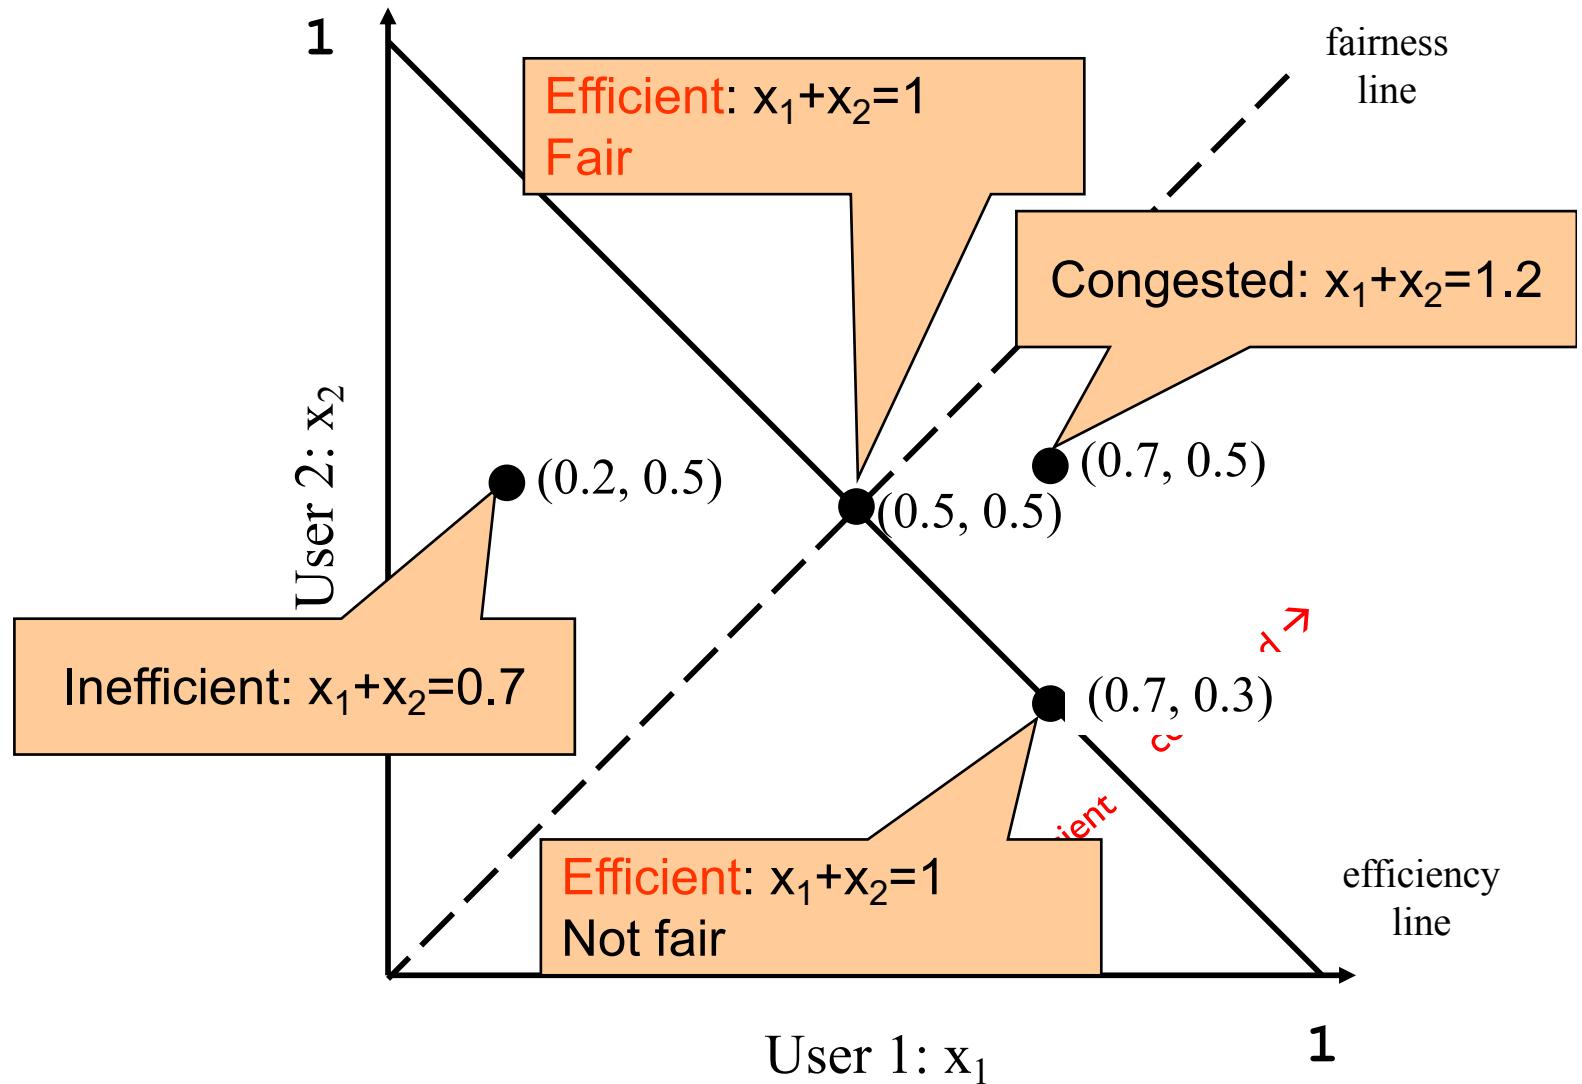
- Some rate adjustment options: Every RTT, we can
 - Multiplicative increase or decrease: $\text{CWND} \rightarrow a * \text{CWND}$
 - Additive increase or decrease: $\text{CWND} \rightarrow \text{CWND} + b$
- Four alternatives:
 - AIAD: gentle increase, gentle decrease
 - AIMD: gentle increase, drastic decrease
 - MIAD: drastic increase, gentle decrease
 - MIMD: drastic increase and decrease

Simple Model of Congestion Control

- Two users
 - rates x_1 and x_2
- Congestion when $x_1+x_2 > 1$
- Unused capacity when $x_1+x_2 < 1$
- Fair when $x_1 = x_2$

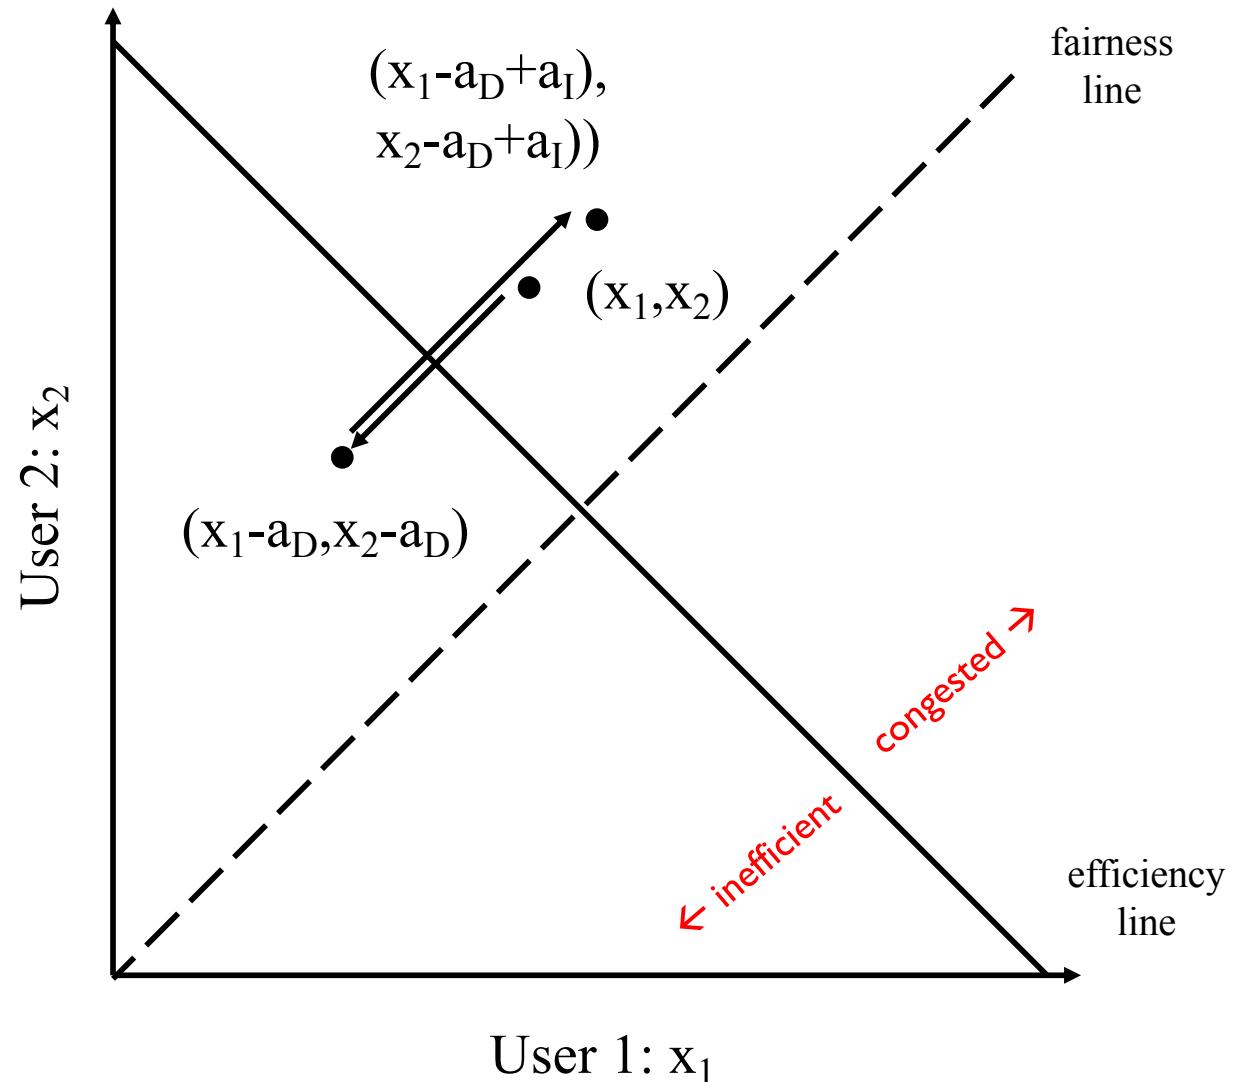


Example

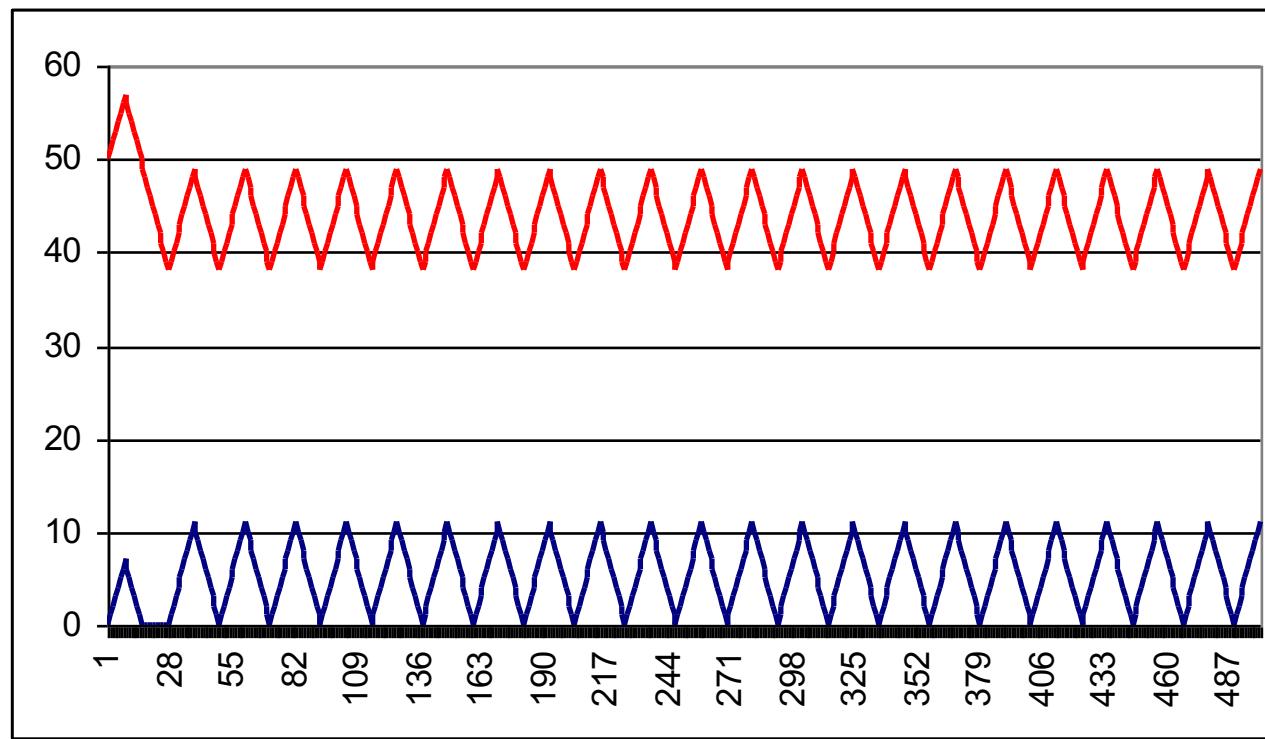
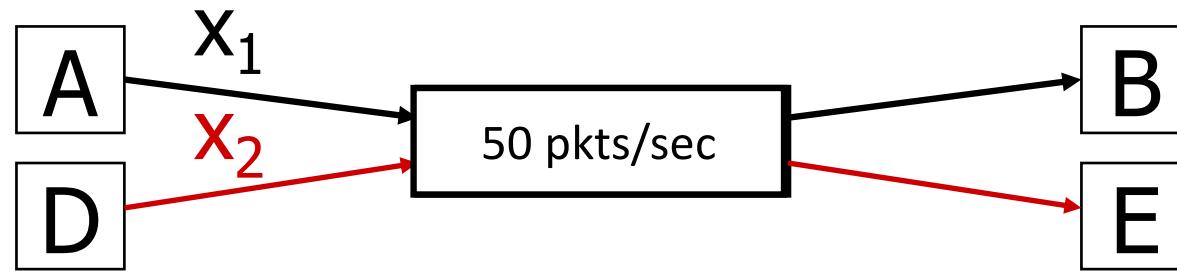


AIAD

- Increase: $x + a_I$
- Decrease: $x - a_D$
- Does not converge to fairness

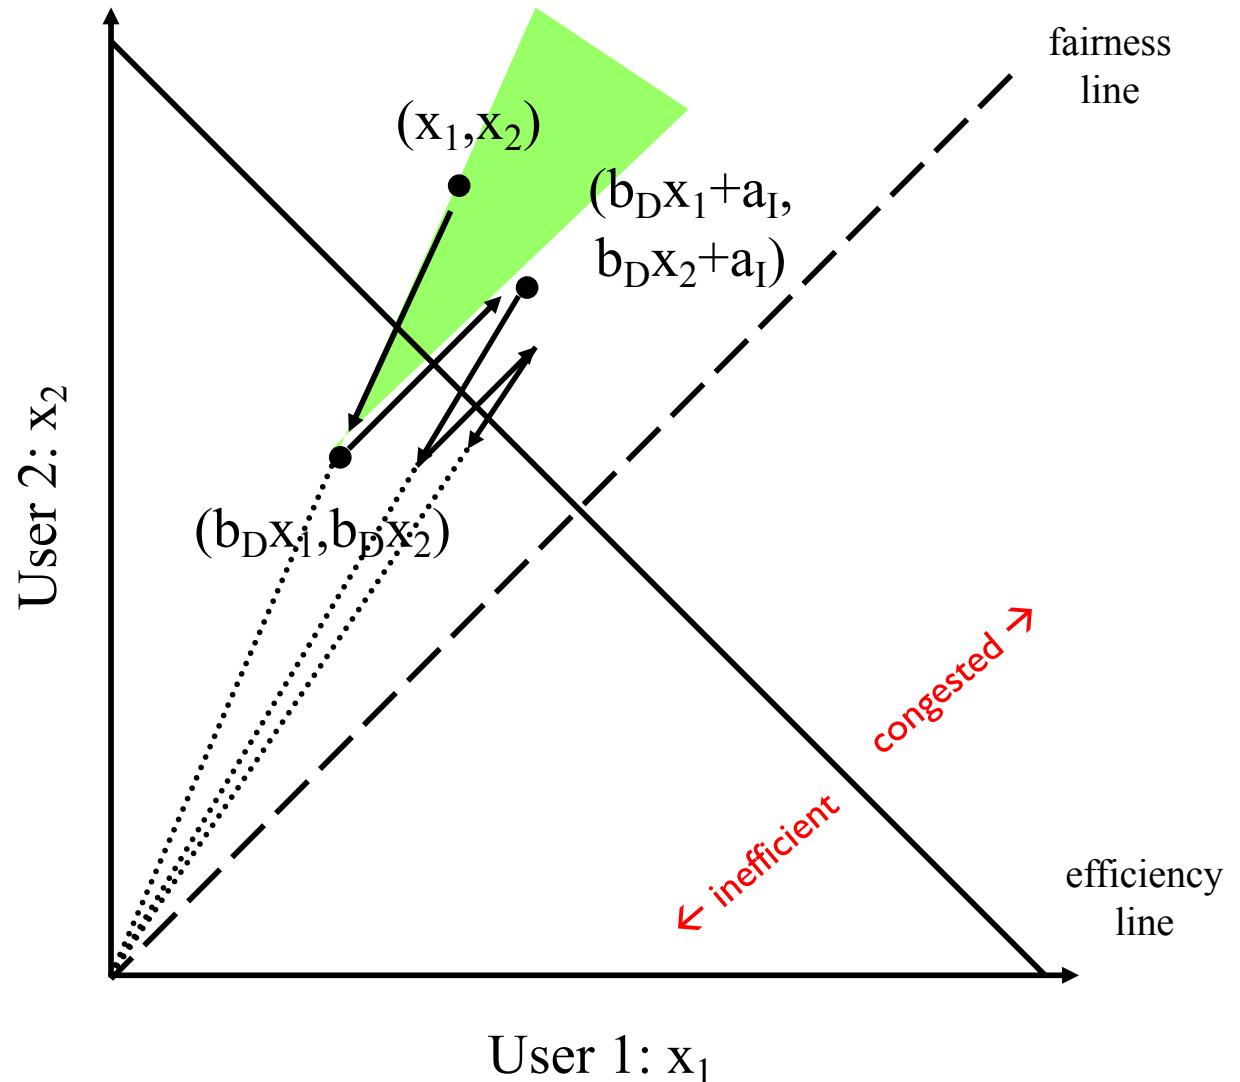


AIAD Sharing Dynamics

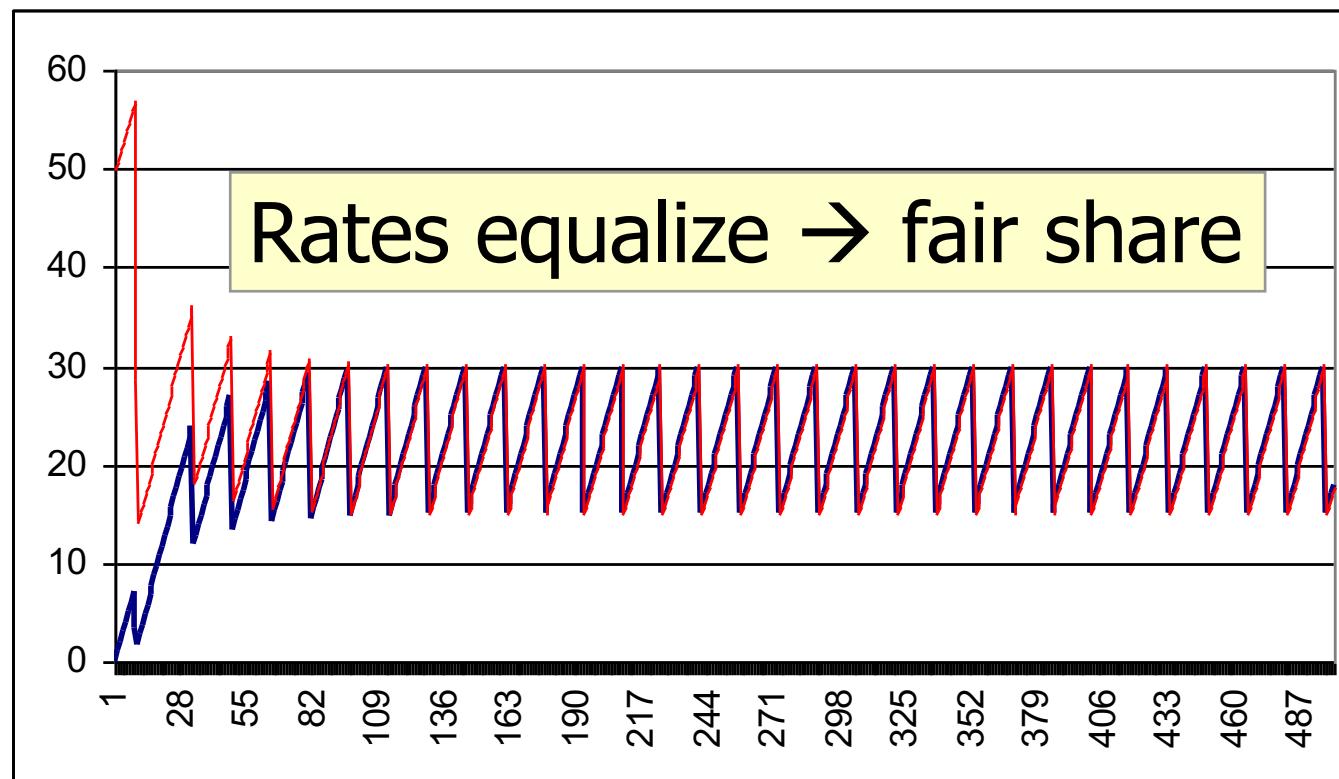
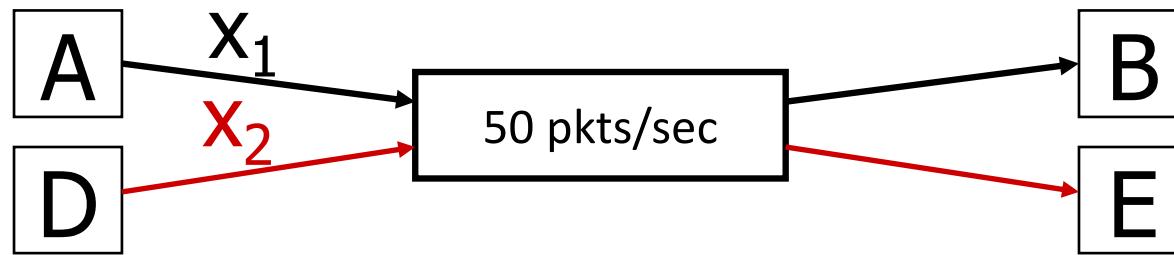


AIMD

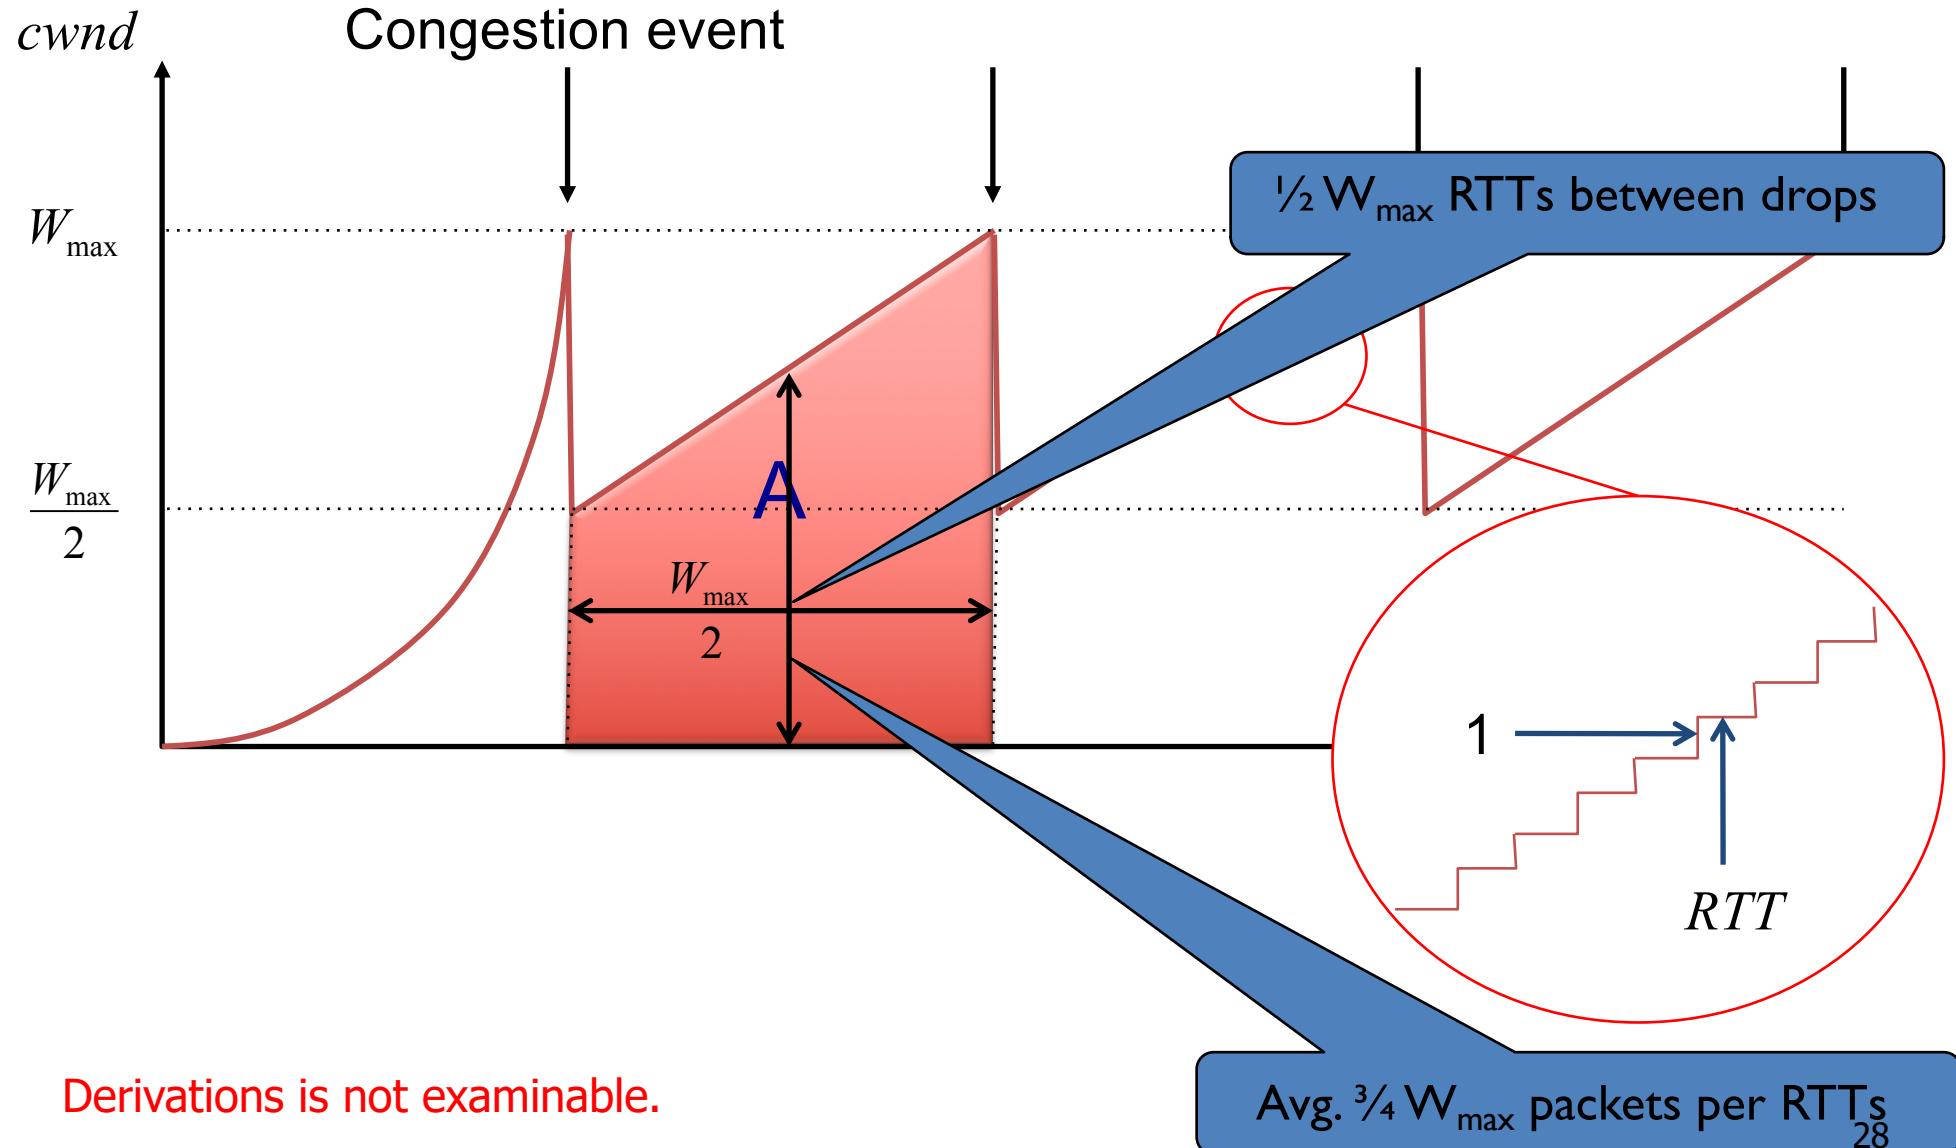
- Increase: $x+a_I$
- Decrease: $x*b_D$
- Converges to fairness



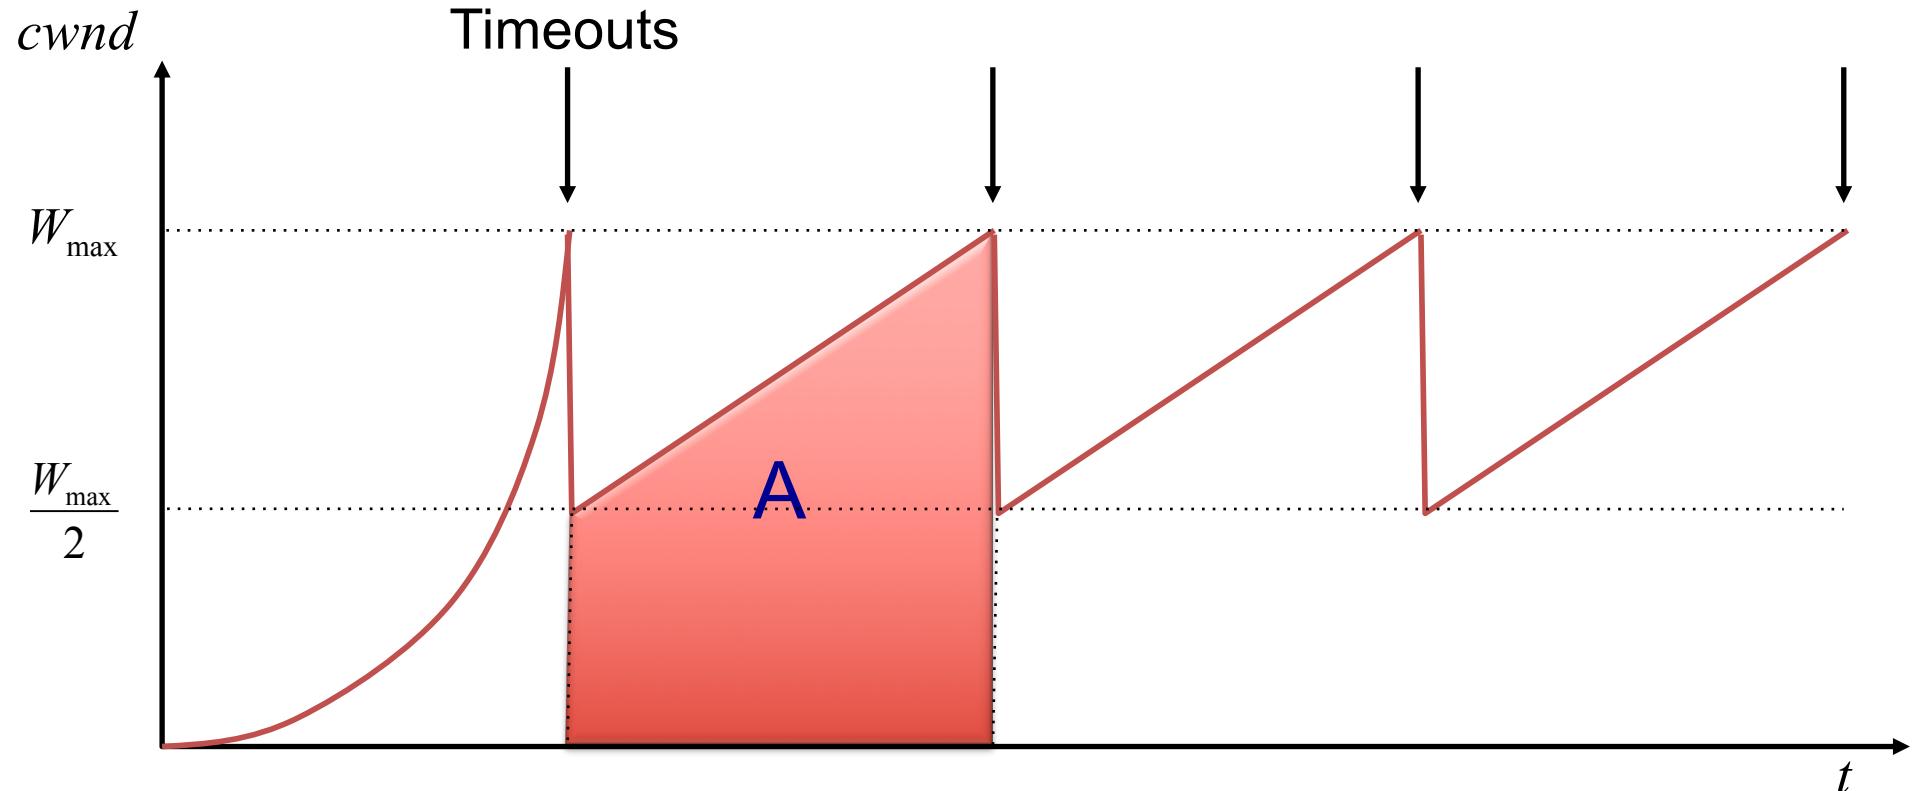
AIMD Sharing Dynamics



A Simple Model for TCP Throughput



A Simple Model for TCP Throughput



Packet drop rate, $p = 1/A$, where $A = \frac{3}{8}W_{\max}^2$

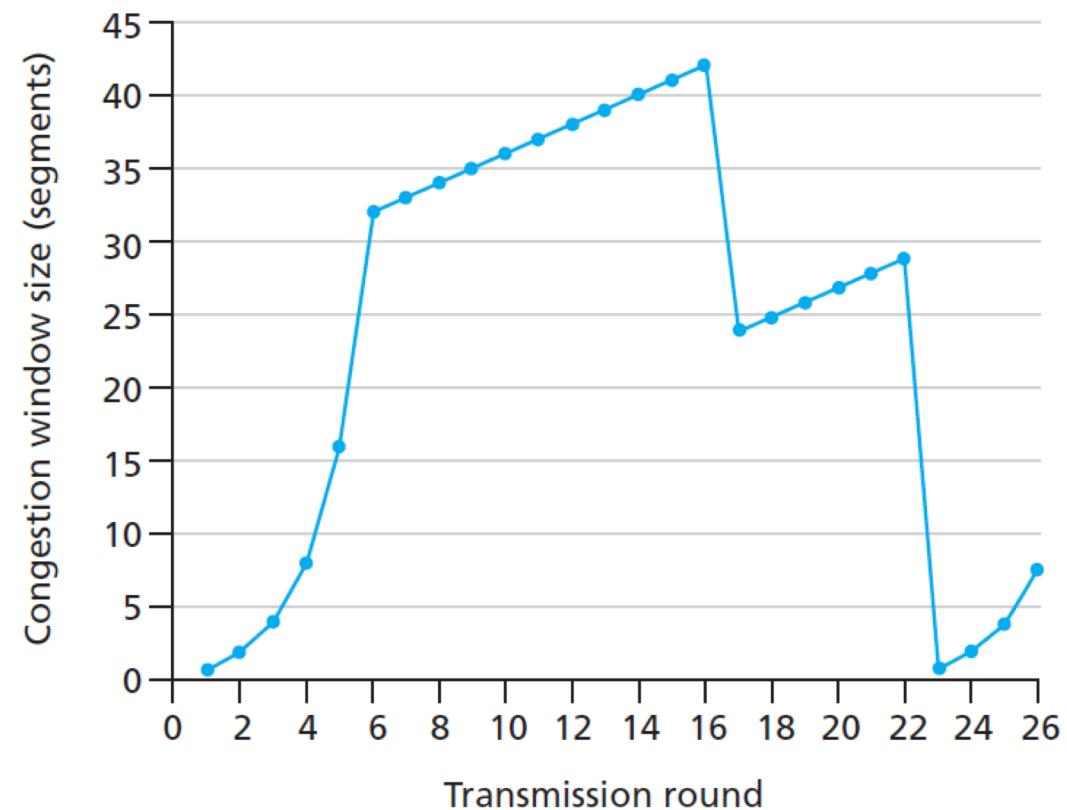
$$\text{Throughput, } B = \frac{A}{\left(\frac{W_{\max}}{2}\right)RTT} = \sqrt{\frac{3}{2}} \frac{1}{RTT \sqrt{p}}$$

Quiz: TCP Congestion Control?



In the figure how many congestion avoidance intervals can you identify?

- A. 0
- B. 1
- C. 2
- D. 3
- E. 4

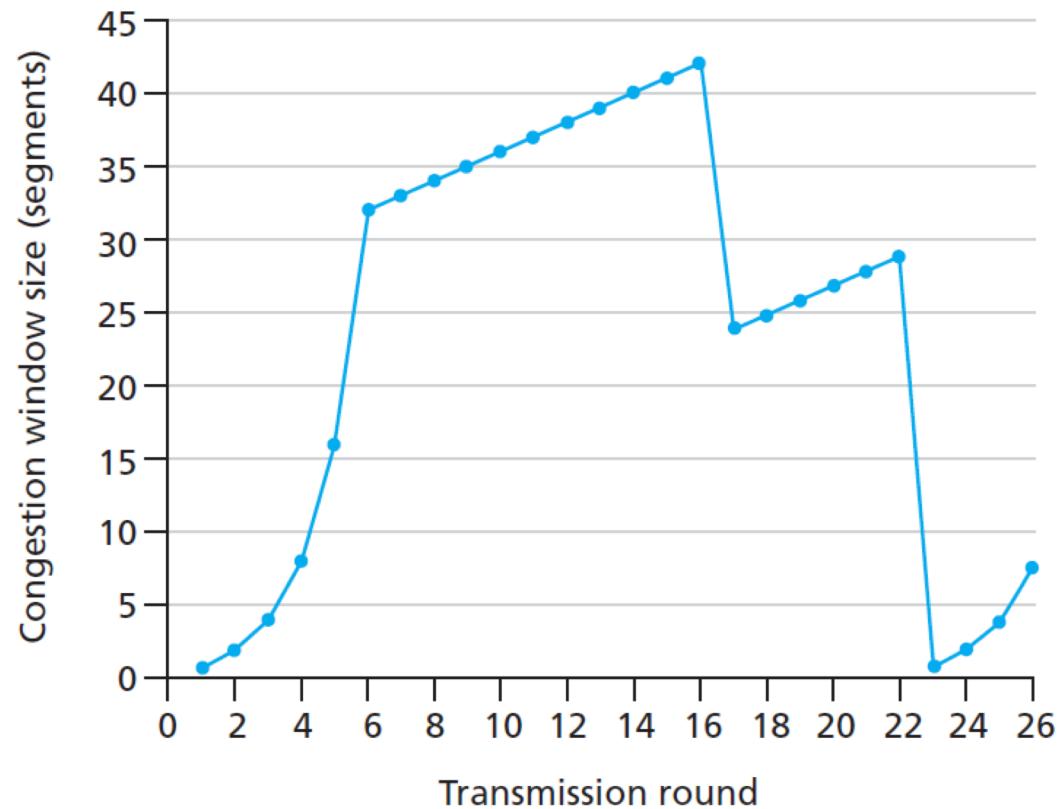


Quiz: TCP Congestion Control?



In the figure how many slow start intervals can you identify?

- A. 0
- B. 1
- C. 2
- D. 3
- E. 4

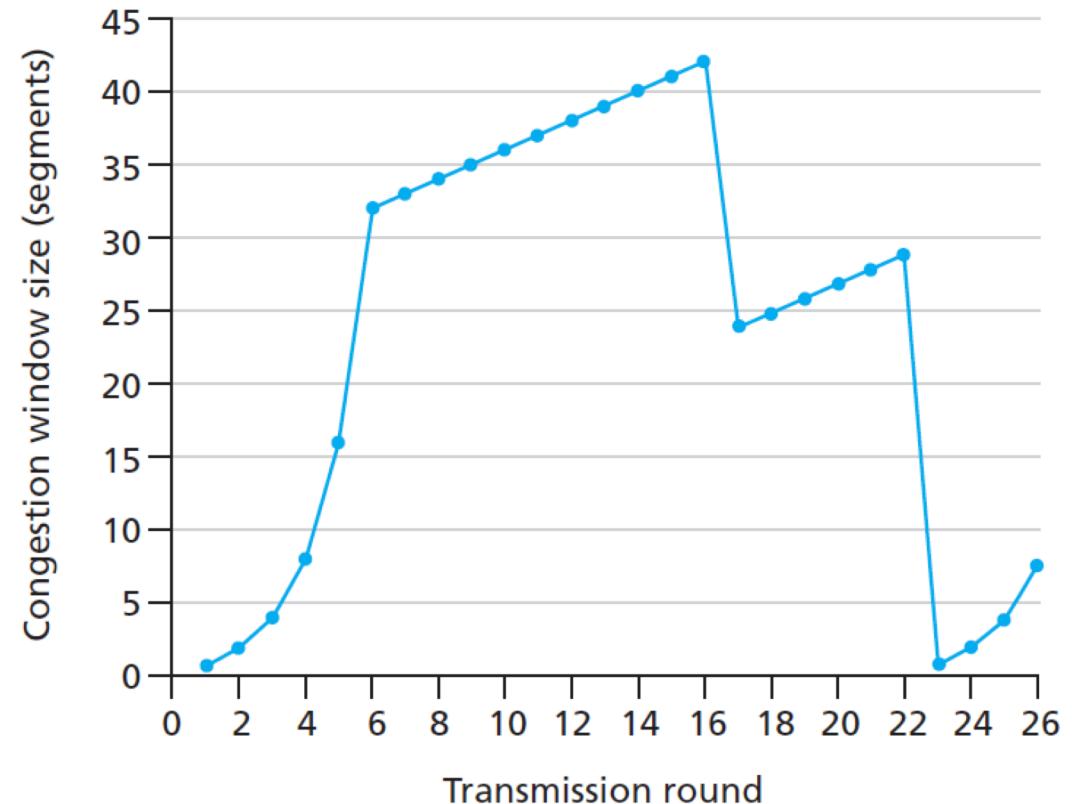


Quiz: TCP Congestion Control?



In the figure after the 16th transmission round, segment loss is detected by ?

- A. Triple Dup Ack
- B. Timeout

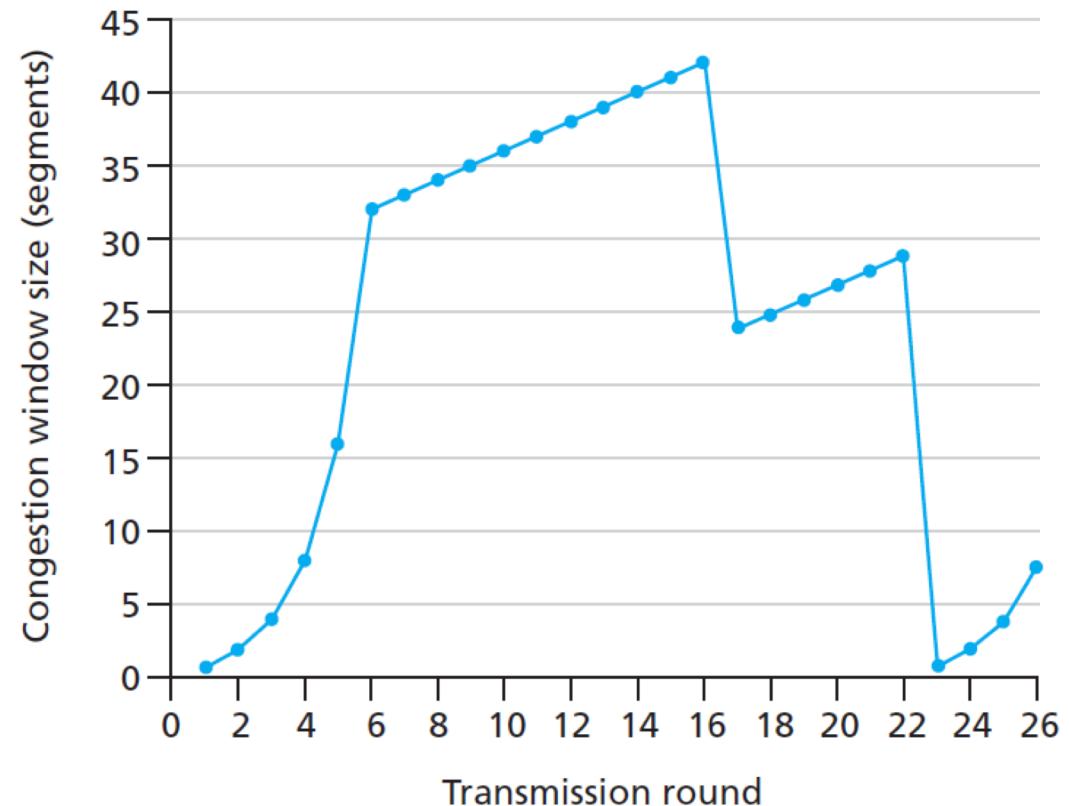


Quiz: TCP Congestion Control?



In the figure what is the initial value of ssthresh (slow start threshold)?

- A. 0
- B. 28
- C. 32
- D. 42
- E. 64



Transport Layer: Summary

- ❖ principles behind transport layer services:
 - multiplexing, demultiplexing
 - reliable data transfer
 - flow control
 - congestion control
- ❖ instantiation, implementation in the Internet
 - UDP
 - TCP

next:

- leaving the network “edge” (application, transport layers)
- into the network “core”